

SENG 440: *Computing Convergence Method*

By: Matthew Trent (V00982038) · Gabriel Maryshev (V00993574)

15 slides excluding 5 benchmark graphs

CCM for the Raspberry Pi 4 B 🖥️

- It's a “shift and add”-type algorithm for calculating transcendental functions.
- Uses *fixed-point*, opposed to *floating-point arithmetic* for efficiency and precision.
- Our implementation is tailored specifically for our *Raspberry Pi 4 B*, 64-bit ARM Cortex-A72 processor, ARMv8-A architecture, gcc-optimized.



My Raspberry Pi 😊

Design requirements

1. We need to select one of four transcendental functions to perform CCM on: $\log_2(M)$, e^M , $M^{(1/2)}$, or $M^{(1/3)}$.
2. We need to take in inputs wider than the ideal theoretical range.
3. CCM needs to be implemented using fixed-point arithmetic.
4. Determine the bottleneck(s) in the CCM algorithm.
5. Optimize the algorithm.
6. Determine the speed-up of the new implementation.

Design choices

- We chose the *base-2 logarithm*.
- Argument range:
 - Practical: $0 < M \leq 2^{16}$.
 - Theoretical: $0 < M \leq 0.5$.
- Barr Group and MISRA C's Embedded C Coding Standards.
- Chose scale factor of 2^{15} .

Bit positions (from right)	Allocated bits	Purpose	Description
31st	1 bit	Number's sign	Indicates if the number is positive or negative.
30th - 15th	16 bits	Integer part of number	Represents the integer portion of the number, capable of storing values from -2^{15} to $2^{15} - 1$ ($-32,768$ to $32,767$).
14th - 0th	15 bits	Fractional part of number	Provides fractional precision using a scale factor of 2^{15} , meaning 15 bits of fractional precision.

32 of possible 64-bit delegation

Our run script for the Pi

Purpose: Simple code transfer, compilation, and execution on Raspberry Pi.

- Used like: `./run.sh /path/to/file.c -some -flags -here`.
- Transfers files and compiles with GCC.
- Generates assembly and performance reports.
- Executes binaries, returns output: `~/asm` and `~/stats`.

Default flags:

- `-mcpu=cortex-a72, -O3, -fno-stack-protector, -fomit-frame-pointer, -lm`.

Core implementation

Calculation of Binary Logarithm – Pseudocode

```
1:                                     ▷  $\log_2 M$  with  $K$  bits of precision
2: for  $i = 0$  to  $K - 1$  do
3:    $LUT(i) = \log_2(1 + 2^{-i})$        ▷ calculate the table with  $\log_2 A_i$ 
4:  $f = 0$ 
5: for  $i = 0$  to  $K - 1$  do
6:    $\mu = M \cdot (1 + 2^{-i})$          ▷ potential multiplication by  $A_i$ 
7:    $\phi = f - LUT(i)$                  ▷ potential addition with  $\log_2 A_i$ 
8:   if  $\mu \leq 1.0$  then
9:      $M = \mu$                        ▷ if product is less than 1 accept iteration,
10:     $f = \phi$                        ▷ otherwise reject it (do nothing)
11: return  $f$ 
```



```
#include <stdio.h>
#include <math.h>

// # of bits of precision
#define K 16

void calculate_lut(double LUT[K]) {
    for (int i = 0; i < K - 1; i++) {
        LUT[i] = log2(1 + pow(2, -i));
    }
}

double log2_CCM(double M) {
    double LUT[K];
    calculate_lut(LUT);

    double f = 0;

    for (int i = 0; i < K - 1; i++) {
        double u = M * (1 + pow(2, -i));
        double phi = f - LUT[i];

        if (u <= 1.0) {
            M = u;
            f = phi;
        }
    }

    return f;
}

int main() {
    double M = 0.6;

    printf("unoptimized log2(%f) = %f\n", M, log2_CCM(M));

    return 0;
}
```

Optimizations: part 1 ⚡

Dynamic to defined LUT:

```
void calculate_lut(int32_t LUT[K]) {
    for (int i = 0; i < K - 1; i++) {
        LUT[i] = (int32_t)(log2(1 + pow(2, -i)) * SCALE_FACTOR);
    }
}
```

```
import math
# generate LUT for embedding in the C program
print(", ".join(map(str, [int(math.log2(1 + math.pow(2, -i)) * (1 << 15)) for i in range(15)])))
```

```
const int32_t LUT[K-1] = {
    32768, 19168, 10548, 5568, 2865, 1454, 732,
    367, 184, 92, 46, 23, 11, 5, 2
};
```

Fixed-point arithmetic:

```
// 2^(K-1) represents our scale => 2^15 = 32768
#define SCALE_FACTOR (1 << (K - 1))
```

```
for (int i = 0; i < K - 1; i++) {
    int32_t u = M + (M >> i);
    int32_t phi = f - LUT[i];

    if (u <= SCALE_FACTOR) {
        M = u;
        f = phi;
    }
}
```

```
int main() {
    double M_real = 0.6;
    // convert to fixed-point notation
    int32_t M_fixed = (int32_t)(M_real * SCALE_FACTOR);

    int32_t result_fixed = log2_CCM(M_fixed);
    printf("ccm log2(%d) = %d\n", M_fixed, result_fixed);

    // revert to floating-point notation
    double result_real = (double)result_fixed / SCALE_FACTOR;

    printf("unoptimized fp ccm log2(%f) = %f\n", M_real, result_real);

    return 0;
}
```

Optimizations: part 2 ⚡

SIMD (NEON):

I converted this, our original LUT:

```
int32_t LUT[K - 1] = {32768, 19168, 10548, 5568, 2865, 1454, 732, 367, 184, 92, 46
```

Into:

```
// defining the LUT arrays separately
int32_t LUT_array1[4] = {32768, 19168, 10548, 5568};
int32_t LUT_array2[4] = {2865, 1454, 732, 367};
int32_t LUT_array3[4] = {184, 92, 46, 23};

// this last "-1" is a space filler, we don't need it
// we just want to fill up all 4 32-bit fields for alignment
int32_t LUT_array4[4] = {11, 5, 2, -1};

// loading LUT into NEON vectors
int32x4_t LUT_vec[4] = {
    vld1q_s32(LUT_array1),
    vld1q_s32(LUT_array2),
    vld1q_s32(LUT_array3),
    vld1q_s32(LUT_array4)
};
```

Input normalization:

```
// at the start, we normalize
int shifts = 0;
while (M >= SCALE_FACTOR)
{
    M >>= 1;
    shifts++;
}

// later denormalizing after the main loop
f += shifts << 15; // K - 1 = 15
```


Optimizations: part 3 ⚡

Loop header:

Here's the initial header:

```
for (int i = 0; i < K - 1; i++) { ... }
```

Here's the improved one:

```
for (register int i = 0; i < K - 2; i += 2) { ... }
```

Register usage, ternary operators, bitwise ops, & loop unrolling:

```
for (register int i = 0; i < K - 2; i += 2)
{
    // unrolled loop portion #1
    register int32_t u1 = M + (M >> i);
    register int32_t LUT_val1 = LUT[i];
    register lteSF1 = u1 <= SCALE_FACTOR;
    M = lteSF1 ? u1 : M;
    f = lteSF1 ? f - LUT_val1 : f;

    // unrolled loop portion #2
    register int32_t u2 = M + (M >> (i+1));
    register int32_t LUT_val2 = LUT[i + 1];
    register lteSF2 = u2 <= SCALE_FACTOR;
    M = (lteSF2) ? u2 : M;
    f = (lteSF2) ? f - LUT_val2 : f;
}
```

Optimizations: part 4 ⚡

- Several new ones*, now our *final version*.
- Overall:
 - *Operator strength reduction.
 - *Reducing function call overheads.
 - *Register keywords.
 - *Locality of variable definitions.
 - *Bitwise operations and comparisons.
 - *Software pipelining.
 - Predicate operations.
 - Loop unrolling.
 - SIMD (NEON).
 - Fixed point arithmetic.
 - Simple asymptotic optimization (analysis).

```
#include <stdio.h>
#include <stdint.h>
#include <arm_neon.h>

#define K 16
#define SCALE_FACTOR (1 << (K - 1))

int main()
{
    // conversion to fixed-point notation
    double real_input = 22;
    register int32_t M = (int32_t)(real_input * SCALE_FACTOR);

    register int32_t f = 0;

    // defining the LUT arrays separately
    int32_t LUT_array1[4] = {32768, 29168, 18548, 5568};
    int32_t LUT_array2[4] = {2865, 1454, 732, 367};
    int32_t LUT_array3[4] = {184, 92, 46, 23};
    int32_t LUT_array4[4] = {11, 5, 2, -1}; // -1 to represent we don't use this place of the array, but have the space (for alignment)

    // load the LUT into NEON vectors
    int32x4_t LUT_vec[4] = {
        vldiq_s32(LUT_array1),
        vldiq_s32(LUT_array2),
        vldiq_s32(LUT_array3),
        vldiq_s32(LUT_array4)};

    // normalization of M to range
    int shifts = 0;
    while (M >= SCALE_FACTOR)
    {
        M >>= 1;
        shifts++;
    }

    for (register int i = 0; !(i & 16); i += 2)
    {
        // NEON to LUT value #1 for unroll 1
        int32_t LUT_val1 = vgetq_lane_s32(LUT_vec[i >> 2], i & 3);

        // #1 unrolled iter
        register int32_t u1 = M + (M >> 1);
        register lteSF1 = u1 <= SCALE_FACTOR;
        M = lteSF1 ? u1 : M;
        f = lteSF1 ? f + LUT_val1 : f;

        // pipelining!
        // ...
        // prepare for the next iteration within the current #2 one
        if (!(i + 2) & 16))
        {
            // NEON to LUT value #2 for unroll 2
            int32_t LUT_val2 = vgetq_lane_s32(LUT_vec[(i + 1) >> 2], (i + 1) & 3);
            register int32_t u2 = M + (M >> (i + 1));
            register lteSF2 = u2 <= SCALE_FACTOR;
            M = lteSF2 ? u2 : M;
            f = lteSF2 ? f + LUT_val2 : f;
        }
    }

    // denormalize the fixed-point value
    f += shifts << 15; // K - 1 = 15
    printf("optimized fp ccm log2(%f) = %f\n", 0.6, (double)f / SCALE_FACTOR);

    return 0;
}
```

Flags

- None for versions 1-4.
- Optimized as much as we could for our *final version*, 5:
- These overrode the ones we defined earlier as defaults in our run script.

- `-mcpu=cortex-a72` : Optimize for Cortex-A72.
- `-O3` : Maximize optimization.
- `-fno-stack-protector` : Disable stack protection.
- `-fomit-frame-pointer` : Omit frame pointer.
- `-march=armv8-a` : Set target architecture.
- `-fprefetch-loop-arrays` : Prefetch loop data.
- `-mtune=cortex-a72` : Fine-tune for Cortex-A72.
- `-ftree-vectorize` : Enable loop vectorization.
- `-funroll-loops` : Unroll loops.

Assembly output into ~/asm 🐙


- Optimized version:
 - Shows fewer branch misses.
 - Less load/store cycles.
 - Smaller Instruction count.
 - Efficiently handles conditional operations without branching.
 - Closest equivalent in ARMv8-A is: *csel*, *csinc*, *csinv*, *cset*, etc.
 - Our code seemed *optimized-enough without diving much into assembly details*.

Benchmarking

Used Linux's Perf to benchmark all 5 versions:


Comparable versions	Version	Page faults	Cycles	Instructions	Branch misses	ASM file length
✗	1_base.c (hyper-optimized library function)	55	379,156	104,739	2,536	33
Initial 🐢	2_unoptimized.c	55	392,134	110,437	2,628	140
✗	3_fixed_point_arithmetic.c	55	378,153	106,663	2,575	349
✗	4_defined_lut.c	46	333,282	85,428	1,984	140
Final 🐇	5_general_optimizations.c	47	326,198	85,037	1,925	33

Page faults




16% less page faults.

Branch misses




31% fewer branch misses.

ASM length




124% fewer lines of ASM.

Cycles



18% fewer cycles.

Instructions



26% fewer instructions.

Proof of correctness

- Ran 1,000 times for certainty.
- Random $0 < M \leq 100$ values.
- Could test up to 2^{16} .
- 0.030550% mean deviation from “*true log₂*”.
- Expected difference given memory and speed trade-off.

... ^ 998 more cases ^ ...

TEST CASE 999:

Randomly chosen input: 34.719630

CCM log2: 5.117615

True log2: 5.117680

Percent difference: 0.001269%

TEST CASE 1000:

Randomly chosen input: 53.251415

CCM log2: 5.734711

True log2: 5.734748

Percent difference: 0.000650%

MEAN OVERALL PERCENTAGE DIFFERENCE: 0.030550%

Conclusions and profiling

- Would have used Valgrind, but:
 - Non-long running.
 - Weren't concerned about memory leaks.
- Would have used Cachegrind, but:
 - That focused on cache optimization while (most of) ours worked on algorithmic speed-up.
- Achieved *43% improvement* on-average across *5 key algorithm metrics*.